AI Audit: FreedAI

Table of Contents

Introduction 2
History 4
Inputs 6
Outputs 8
Methodology 9
    Results of Experiment 11
Consequences 14
Works Cited 16

Introduction

A year ago, I contracted a serious gastrointestinal illness while abroad in Japan. I was almost entirely bed-bound for three days, resigned to a brutal cycle of rehydrating the fluids I was rapidly losing with strange electrolyte jelly drinks in a hotel room only marginally larger than the mattress which laid at its center. My condition improved upon my return home, but slowly. On the same day I returned, I visited the nearest urgent care clinic, desperate for a solution. Finally face-to-face with a professional who could diagnose my condition, my doctor asked a bizarre question: "Is it okay if I record our conversation so my AI can take notes?"

I was baffled. In my mind raced several questions. Is this a violation of HIPPA? How will my doctor use AI to diagnose my illness? What program is she using? Why is she using a note-taking AI at all? I wanted to say *no, I don't want you to use AI,* but I was so desperate for an answer to this painful illness, I didn't want to say no and risk annoying my doctor in any way. I was caught in a perplexing situation: say no and risk upsetting my doctor, the person who held answers and treatment to my illness; or say yes and allow my doctor to resume without indisposing her ability to treat me. I said yes, the AI could record me.

This interaction still weighs uncomfortably on me. We still understand very little about how user privacy is maintained by AI organizations, if at all (Lindgren et al. 2025). Who is to say whatever program my doctor used (I was not informed of its name) properly removes my medical information from its data center, or does not save my voice for internal training (Kong & Ding 2024)? It is unclear if I was given the proper resources to make an informed, consensual decision in this situation, as this technology is being adopted in healthcare so quickly that ethnical adaptations are lagging behind. As a result, I, the patient, was left vulnerable, fearing my inability to consent would interfere with the quality of care I received.

To demystify the medical note-taking process I was not privy to in this situation, I decided to research what AI programs clinicians such as my urgent care doctor are using. She might have used a medical scribe AI, an emerging technology "that listen[s] to physician-patient interaction and generate[s] an impressively detailed subjective, objective, assessment, and plan (SOAP) note, together with an email summary intended for the patient" (Mess et al. 2025). Medical scribe AI programs are endorsed and funded by organizations such as the American Medical Association and The Permanente Medical Group and popularized by drastic claims to improve clinicians' quality of life; to "restore the human side of medicine… reducing physician burnout by removing administrative burdens and providing real-world solutions to help doctors rediscover the Joy in Medicine" (Feldheim 2025).

In this audit, I identify a popular representative of medical scribe AI technology, Freed AI, and replicate my interactions at the urgent care clinic to unveil how clinicians interact with medical scribe AIs. I weigh the experience of patients who are unable to engage directly with the AI against the clinician's, noting possible violations of patient privacy and spaces for ethical patient-clinician-AI interactions.

History

A medical scribe is “a nonclinician who assists with documentation in the medical record on behalf of the treating clinician. The scribe may also assist with other work flow components (eg, order entry, examination room preparation) in certain settings” (Gottlieb 2020). Since the introduction of electronic health records (EHRs) in the 1960s, medical scribes have used digital technology to document and maintain information about patients, allowing physicians to focus more on patient care than administration (Atherton 2011).

AI medical scribes emerged due to two crucial advancements in technology: public access to large language models (LLMs) and voice recognition technology. Before these technologies converged, machine learning systems lacked the ability to interpret dialogue and synthesize it into context-sensitive text.

It is difficult to determine any one AI medical scribe that originated first because many early tools may have been proprietary or unpublished, making them hard to trace historically. However, an EHR technology known as contextual autocomplete was described by doctors in 2020, able to predict inputs as a clinician writes them (Agrawal et al. 2020). It did not fully scribe conversations, but it helped automate part of the note-creation process, creating a precedent for AI medical scribes. So, while there were earlier machine-assisted systems in medicine, the first AI medical scribe emerged in 2023-2024 as industries worldwide adopted LLMs (Feldheim 2025).

The history of Freed AI is more traceable though limited. Co-founder and CEO Erez Druk describes creating Freed to help his wife, a clinician who often worked late hours to write patient documentation, and allow her to enjoy free time with family (Druk 2024). Freed is explicitly clinician-first, designed to reduce burnout, increase efficiency, and return time to medical professionals (Druk 2024; Chen and Lee 2025). The target audience is, undeniably, clinicians.

While some clinicians praise AI scribes for increasing the efficiency of their patient interactions, many others critique them. “The notes it produces are verbose and formulaic. Lots of information, not enough clarity and meaning. I find myself having to insert my deliberations: how and why I came to the conclusions I did; why I decided on one course of action and discounted other possible options” (Whitaker 2025). Even with the rapid speed which these notes are made, medical professionals must validate the notes, a task which grows more laborious the more hallucinations occur. Still, the technology continues to appear in EHRs, and as they are so do concerns.

Inputs

As AI medical scribes rely on LLMs, so too do they rely on data centers, including the land, fresh water, electricity, minerals, and metals these facilities consume. Similarly, the biases imprinted by institutions and coders onto LLMs are also present in Freed, meaning responses may be unfairly swayed by prejudices of race, gender, and ability (Lindgren et al. 2025; Kong and Ding 2024).

The Freed program itself relies on transcripts of audio conversations (Figure 1).

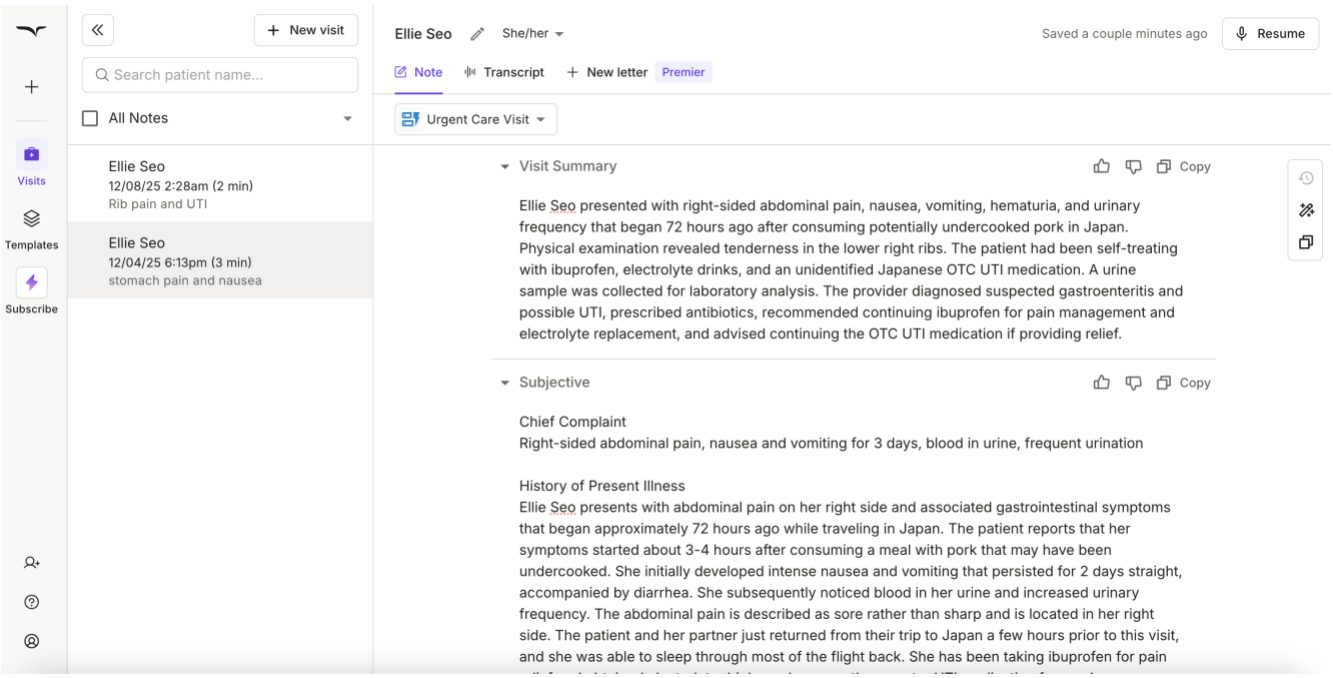

Figure 2: Freed AI homepage.

In practice, this means that every spoken word (including symptoms, speculations, emotional expressions, hesitations, side comments, and clinician reasonings) are captured, converted into text, and fed into a LLM that synthesizes the information into a clinical note. This clinical note contains eight distinct sections:

1. Transcript: contains the Summary and Full Transcript of the conversation, though both read only as transcripts. The Summary contains a paragraph break between the speech of

each person whereas the Full Transcript contains a paragraph break after every sentence. It seems the only difference is formatting.

2. Visit Summary: details of the conversation condensed in paragraph form.
3. Subjective: details gathered from patient, including their chief compliant, pre-existing conditions, and social history.
4. Objective: details gathered from clinician.
5. Assessment & Plan: summary of patient visit and symptoms including final determinations made by clinician.
6. Clinical Codes: alphanumeric values that quickly communicate medical information.
7. Patient Instructions: letter to be sent to the patient containing conclusions from the visit.

Outputs

The main product of Freed is documentation for clinicians. Structured notes are taken from the transcript, including summaries and conclusions. Transcripts and clinical codes are also generated by Freed. These results intend to replicate the note-taking process for clinicians so they don't need to, reducing the time spent on administrative work and therefore the risk of burnout.

A great risk of AI medical scribes is in the hallucinations that may be present in any LLM technology (Mess et al. 2025). The AI may fabricate symptoms, instructions, or clinical details not present in the conversation. The severity of even small hallucinations becomes very serious when institutionalized in medical records.

Additionally, the quality of writing produced by AI medical scribes has been described as "verbose and formulaic," requiring constant interjections for best usage (Whitaker 2025). In sum, it is impossible to produce ethical, labor-free notes with an AI medical scribe like Freed.

Methodology

To investigate how a clinician might interact with Freed AI during a real medical encounter, I designed a simulation which replicates the actual interactions I had with my urgent care physician a year ago. Because access to Freed was limited to a seven-day trial window, I developed a plan prior to ensure that all stages of the experiment could be completed efficiently.

My aim in replicating this event was to provide a reference point I could use to then explore how Freed processes patient narratives, clinician reasoning, and conversational context. This design also served a theoretical purpose: to undo the black box surrounding my original visit and illuminate how my clinician may have interacted with AI tools.

I drafted a script to be read in two parts: one for the first visit and a second for the follow-up visit (Figure 2). The patient was named Ellie Seo and the clinician Syndey Alistar. This script was copied into Typecast, a program which reads text aloud using AI voices. Freed listened to these conversations.

INT. URGENT CARE CLINIC - EXAM ROOM

DOCTOR
Okay, I've started recording. So, what brings you in today?

PATIENT
I've been experiencing pain in my stomach, around my right side. My partner and I just got back from a trip to Japan a few hours ago, and about three days ago I started having really intense nausea. I was throwing up for two days straight, and then I noticed blood in my urine. I needed to pee frequently. I was lucky to sleep through most of the flight back.

DOCTOR
That sounds uncomfortable. So your symptoms began about 72 hours ago?

PATIENT
Yes.

DOCTOR
Do you have any idea what might have caused your symptoms?

PATIENT
I think it's food poisoning. I had a meal with pork that might have been uncooked. About three or four hours after I had that, I started vomiting and having diarrhea.

DOCTOR
That could be it. You might have bruised yourself from vomiting, too, which could explain the pain in your side. Is it a sharp pain or a dull pain?

PATIENT
Dull.

DOCTOR
Okay. Is it okay if I begin your physical exam?

PATIENT
Yes.

DOCTOR
Okay.

DOCTOR approaches PATIENT on the exam chair.

DOCTOR
I'm going to press around your right side. Does it hurt when I do this?

PATIENT
Not much with pressure.

DOCTOR
Okay. How about here, ribs?

PATIENT
Oof, yeah.

DOCTOR
Okay. Lower ribs on right side are tender. Well, we'll run labs on the urine sample you gave to see if you have a UTI. The discomfort in your lower side area could be from that, too. Have you taken any medication for the pain?

PATIENT
We visited a pharmacy right before we left. I had a lot of electrolyte drinks and had a UTI over-the-counter medicine, but I don't know what's in it because it was all in Japanese (laughs). It's been helping, I think. I've also been taking ibuprofen.

DOCTOR
Okay. I'll send you home with some antibiotics. We want to begin treatment if it's something more serious like gastroenteritis. You can also take some UTI over-the-counter medicine if it brings you relief and ibuprofen. We'll have the lab results in a couple days or so, which you can see on your account. Other than that, you're free to go.

PATIENT
Great, thank you.

DOCTOR
Take care.

END FIRST SCENE

INT. URGENT CARE CLINIC - EXAM ROOM

DOCTOR
Okay, I've started recording. What brings you back? How are your symptoms progressing?

PATIENT
The pain in my side and ribs is still here with not much difference. It's like a 7/10, dull pain.

DOCTOR
And we found that you have a UTI, correct?

PATIENT
Yes, the lab work said there was blood in the urine. My discomfort is going away, though.

DOCTOR
Good. Well, I'm going to begin your physical exam. Does it hurt here?

PATIENT
No.

DOCTOR
Pressure on your ribs.

PATIENT
Ouch. Yup.

DOCTOR
What about side, here?

PATIENT
Yeah, there too.

DOCTOR
Well, the antibiotics didn't change your reaction. I think you might have just bruised your ribs when you were vomiting.

PATIENT
Oh, really? So there's no infection.

DOCTOR
I don't think so. Keep taking the antibiotics just in case, but I'm going to prescribe anti-inflammation and muscle relaxants which should reduce your soreness.

PATIENT
Oh, that's a relief. That makes so much sense, now.

DOCTOR
Would you like to try on a brace? It applies pressure which might bring some relief.

PATIENT
Sure, that sounds great. Thank you.

DOCTOR
Of course. I'll have the nurses bring it in shortly.

Figure 2: Script read aloud to Freed AI.

Results of Experiment

For each fictional visit, Freed produced 11-12 passages varying from 1-8 sentences each. These passages had concise prose and borrowed terminology spoken by the participants. For example, in the second visit summary under the Subjective section, Freed writes,

> Ellie Seo returns for follow-up of ongoing side and rib pain. The patient reports persistent pain in her side and ribs with no significant improvement, describing it as a 7/10 dull pain. She states that her discomfort is going away, though the rib and side pain persists. The patient has a diagnosed UTI for which she is currently taking antibiotics.

Freed writes from an observer's perspective, narrating in present-tense as "the patient" describes symptoms almost exactly as the script did, particularly the "7/10 dull pain." It packs details into a singular paragraph for accessible reading by medical professionals.

Freed generated its own headings to organize content, particularly because of the large amount of complex information that must be recorded. Occasionally, Freed used bullet list items to list symptoms, medications, patient histories, and directions.

Freed succeeded in condensing these 2-minute conversations into succinct texts. However, it produced a singular hallucination on the second visit: clinical code S22.9XXA, which describes a fracture of the bony thorax which did not occur (Figure 3). Had an incorrect code like this gone unnoticed in a live clinical setting, the patients medical record risks being marked with a serious physical condition that is unfounded, resulting in potential medical malpractice.

Clinical Codes Premier

Copy

ICD-10 codes

S22.9XXA	Fracture of bony thorax, part unspecified, initial encounter for closed fracture
N39.0	Urinary tract infection, site not specified
R31.9	Hematuria, unspecified

+ Add code

Figure 3: Clinical codes screenshot from simulation of the first visit.

In each simulated interaction, Freed also produced a letter to the patient which summarizes any treatment instructions. Each letter followed the same pattern, opening with a salutation, bullet list items, then inviting questions or concerns (Figure 4).

Patient Instructions Premier

Thank you for visiting today. Here is a summary of the key instructions:

Medications:
- Take antibiotics as prescribed to treat possible gastroenteritis
- You may continue taking ibuprofen for pain relief
- You may continue taking over-the-counter UTI medicine if it helps with your symptoms

Lab Tests:
- Urine test results will be available in a couple days
- Check your patient account online to view the lab results

Disposition:
- You are discharged home today

Please reach out if you have any questions or concerns.

Best Regards,

Sydney Alistar, MD
Urgent Care

Patient Instructions Premier

Dear Ellie Seo,

Thank you for visiting today. Here is a summary of the key instructions:

Medications:
- Continue taking the antibiotics as prescribed

- Take the new anti-inflammatory medication to reduce soreness

- Take the new muscle relaxants to reduce soreness

Medical Devices:
- Use the rib brace that will be provided by the nurses

Please reach out if you have any questions or concerns.

Best Regards,

Sydney Alistar, MD
Urgent Care

Figure 4: Patient instructions for first visit (left) and second visit (right).

The style of each letter is precise and, though justly lacking filler, may be considered curt in its brevity. In the first visit's letter, the declaration that "You are discharged home today" is blunt and obvious, making it seem that no treatment had been decided. The discrepancy in spacing between the two letters is strange as well.

Overall, this experiment proved that, while hallucinations do occur in Freed, the program is capable of rapidly producing text from the perspective of an objective, informed medical professional. The errors and (when appropriate) tone must be revised for full efficacy, but the ability of Freed as a note-taking and transcribing service is evident.

Consequences

The deployment of AI medical scribes like Freed produces consequences that extend beyond efficiency gains or improved EHR documentation. Understanding these effects requires asking who benefits, who suffers, and how the boons and banes of the system are distributed across patients, clinicians, and institutions.

One of the most immediate consequences concerns patient consent. Although Freed is designed to serve clinicians, its inputs rely on patient speech, including sensitive, emotional, or private disclosures. Patients may feel obliged to give their consent—as I was—and not fully aware that their words are being fed into a black box algorithm. This asymmetry of power raises ethical questions: Who understands the AI system, and who is acted upon by it? Patients bear the privacy risks, yet the clinician solely benefits.

Another consequence lies in the perceived authenticity of care. The presence of an AI intermediary can make the patient-clinician interaction feel mediated or transactional. For some patients, knowing that an AI is generating the clinician's notes—perhaps even shaping the structure of the conversation—may signal less personalized attention where the AI becomes a silent third participant in the room. This subtle shift may erode trust, especially among patients already wary of institutional medicine.

The benefits of Freed are undeniably the rapid generation of text, accurate voice transcription, and organized interface. However, the medical notes must be screened by the clinician for accuracy, especially because my experimentation revealed that Freed is subject to hallucinating serious medical conditions. Freed also struggles to write in a polite, consistent manner for patients, as seen in both Patient Instructions forms the program generated. The

involvement of Freed in a clinical setting might mean patients indirectly benefit from a more focused clinician, but the advantages overwhelmingly align with provider interests.

Ultimately, Freed's benefits align primarily with clinicians and health systems. Its burdens—privacy concerns, reduced authenticity, and hallucinated medical details—fall most heavily on patients, especially those with less power to contest or question the records created about them. The consequences of Freed are not simply technical but socio-political. They reveal how AI systems in healthcare often redistribute labor upward while redistributing risk downward. Any meaningful evaluation of Freed must therefore consider not only what it produces, but for whom, and at what cost.

Works Cited

Agrawal, M., Gopinath, D., Murray, L., Horng, S., Karger, D., & Sontag, D. (2020). Fast, Structured Clinical Documentation via Contextual Autocomplete. *Machine Learning for Healthcare.* https://doi.org/10.48550/arXiv.2007.15153

Atherton, J. (2011). Development of the Electronic Health Record. *AMA Journal of Ethics*, *13*(3), 186–189. https://doi.org/10.1001/virtualmentor.2011.13.3.mhst1-1103

Chen, J., & Lee, J. (2025, March 5). Partnering with Freed: An AI-Powered Clinician's Assistant. *Sequoia Capital*. https://sequoiacap.com/article/partnering-with-freed-an-ai-powered-clinicians-assistant/

Erez, D. (2024, November 1). *Freeding Clinicians: Our Founding Story*. https://www.getfreed.ai/blog/freed-founding-story

Feldheim, B. (2025, June 12). *AI scribes save 15,000 hours—And restore the human side of medicine*. American Medical Association. https://www.ama-assn.org/practice-management/digital-health/ai-scribes-save-15000-hours-and-restore-human-side-medicine

Gottlieb, M., Palter, J., Westrick, J., & Peksa, G. D. (2021). Effect of Medical Scribes on Throughput, Revenue, and Patient and Provider Satisfaction: A Systematic Review and Meta-analysis. *Annals of Emergency Medicine*, *77*(2), 180–189. https://doi.org/10.1016/j.annemergmed.2020.07.031

Kong, Y., & Ding, H. (2024). Tools, Potential, and Pitfalls of Social Media Screening: Social Profiling in the Era of AI-Assisted Recruiting. *Journal of Business and Technical Communication*, *38*(1), 33–65. https://doi.org/10.1177/10506519231199478

Lindgren, C. A., Yunes, E., & Itchuaqiyaq, C. U. (2025). Technical Communication's Fight

Against Extractive Large Language Modeling by Applying FAIR and CARE Principles of Data. *Journal of Business and Technical Communication*, *39*(1), 11–25. https://doi.org/10.1177/10506519241280587

Mess, S. A., Mackey, A. J., & Yarowsky, D. E. (2025). Artificial Intelligence Scribe and Large Language Model Technology in Healthcare Documentation: Advantages, Limitations, and Recommendations. *Plastic and Reconstructive Surgery Global Open*, *13*(1). https://doi.org/10.1097/GOX.0000000000006450

Whitaker, P. (2025). Doctors, beware the AI scribe. *New Statesman*. https://www.newstatesman.com/politics/health/2025/07/doctors-beware-the-ai-scribe